

Dynamic Web Development

Lecture 2 – The HTTP Protocol

Hypertext Transfer Protocol

A protocol for conveying hypertext pages across a network.

Invented by Tim Berners-Lee.

Development coordinated by W3C and IETF.

Latest standard is: HTTP v2.0,
defined in: RFC 7540 (May 2015)

A request / response protocol between clients and servers.

The client making the HTTP request is called the user agent. (Browser, spider, robot)

HTTP does not require TCP/IP. It can be used on top of any protocol which provides a reliable transport mechanism.

The HTTP Request

User types a URL into their browser. eg
`http://www.example.com/index.html`

An HTTP client initiates a request by establishing a TCP connection to a specified port on the host (Port 80 is the default).

Server listening on that port waits for an HTTP request message.

Browser sends the following request:

```
GET /index.html HTTP/1.1
```

```
Host: www.example.com
```

```
User-Agent: Mozilla/5.0 (Linux; X11; UTF-8)
```

```
Accept-Charset: UTF-8
```

(Untitled) - Wireshark

File Edit View Go Capture Analyze Statistics Help

Filter: Expression... Clear Apply

| No. | Time | Source | Destination | Protocol | Info |
|-----|----------|---------------|---------------|----------|---|
| 1 | 0.000000 | 192.168.7.2 | 66.249.93.104 | TCP | 1064 > http [FIN, ACK] Seq=0 Ack=0 win=65238 [TCP CHECKSUM INCORRECT] Len=0 |
| 2 | 0.044171 | 66.249.93.104 | 192.168.7.2 | TCP | http > 1064 [ACK] Seq=0 Ack=1 win=6432 Len=0 |
| 3 | 7.103235 | 192.168.7.2 | 171.64.64.183 | TCP | 1065 > http [SYN] Seq=0 Len=0 MSS=1460 |
| 4 | 7.277110 | 171.64.64.183 | 192.168.7.2 | TCP | http > 1065 [SYN, ACK] Seq=0 Ack=1 win=49420 Len=0 MSS=1412 |
| 5 | 7.277150 | 192.168.7.2 | 171.64.64.183 | TCP | 1065 > http [ACK] Seq=1 Ack=1 win=65535 [TCP CHECKSUM INCORRECT] Len=0 |
| 6 | 7.279294 | 192.168.7.2 | 171.64.64.183 | HTTP | GET /~knuth/ HTTP/1.1 |
| 7 | 7.461827 | 171.64.64.183 | 192.168.7.2 | TCP | http > 1065 [ACK] Seq=1 Ack=429 win=48992 Len=0 |
| 8 | 7.467966 | 171.64.64.183 | 192.168.7.2 | TCP | [TCP Previous segment lost] [TCP segment of a reassembled PDU] |
| 9 | 7.467992 | 192.168.7.2 | 171.64.64.183 | TCP | [TCP dup ACK 6#1] 1065 > http [ACK] Seq=429 Ack=1 win=65535 [TCP CHECKSUM INCORRECT] |
| 10 | 7.469417 | 171.64.64.183 | 192.168.7.2 | TCP | [TCP out-of-order] [TCP segment of a reassembled PDU] |
| 11 | 7.469485 | 192.168.7.2 | 171.64.64.183 | TCP | 1065 > http [ACK] Seq=429 Ack=1773 win=63763 [TCP CHECKSUM INCORRECT] Len=0 |
| 12 | 7.469501 | 192.168.7.2 | 171.64.64.183 | TCP | [TCP window update] 1065 > http [ACK] Seq=429 Ack=1773 win=65535 [TCP CHECKSUM INCORRECT] |
| 13 | 7.470164 | 171.64.64.183 | 192.168.7.2 | TCP | [TCP Previous segment lost] [TCP segment of a reassembled PDU] |
| 14 | 7.470186 | 192.168.7.2 | 171.64.64.183 | TCP | [TCP dup ACK 12#1] 1065 > http [ACK] Seq=429 Ack=1773 win=65535 [TCP CHECKSUM INCORRECT] |
| 15 | 7.504532 | 192.168.7.2 | 171.64.64.183 | TCP | 1066 > http [SYN] Seq=0 Len=0 MSS=1460 |
| 16 | 7.679444 | 171.64.64.183 | 192.168.7.2 | TCP | http > 1066 [SYN, ACK] Seq=0 Ack=1 win=49420 Len=0 MSS=1412 |
| 17 | 7.679483 | 192.168.7.2 | 171.64.64.183 | TCP | 1066 > http [ACK] Seq=1 Ack=1 win=65535 [TCP CHECKSUM INCORRECT] Len=0 |
| 18 | 7.681699 | 192.168.7.2 | 171.64.64.183 | HTTP | GET /~knuth/diamondsigns/style.css HTTP/1.1 |
| 19 | 7.863225 | 171.64.64.183 | 192.168.7.2 | TCP | http > 1066 [ACK] Seq=1 Ack=423 win=48998 Len=0 |
| 20 | 7.866435 | 171.64.64.183 | 192.168.7.2 | TCP | [TCP Previous segment lost] [TCP segment of a reassembled PDU] |
| 21 | 7.866458 | 192.168.7.2 | 171.64.64.183 | TCP | [TCP dup ACK 18#1] 1066 > http [ACK] Seq=423 Ack=1 win=65535 [TCP CHECKSUM INCORRECT] |
| 22 | 7.867333 | 171.64.64.183 | 192.168.7.2 | TCP | [TCP out-of-order] [TCP segment of a reassembled PDU] |
| 23 | 7.867373 | 192.168.7.2 | 171.64.64.183 | TCP | 1066 > http [ACK] Seq=423 Ack=620 win=64916 [TCP CHECKSUM INCORRECT] Len=0 |
| 24 | 7.874202 | 192.168.7.2 | 171.64.64.183 | HTTP | GET /~knuth/don.gif HTTP/1.1 |

Frame 6 (482 bytes on wire, 482 bytes captured)

Ethernet II, Src: Gigabyte_80:6e:ae (00:1a:4d:80:6e:ae), Dst: Turbonet_d5:60:0d (00:30:eb:d5:60:0d)

Internet Protocol, Src: 192.168.7.2 (192.168.7.2), Dst: 171.64.64.183 (171.64.64.183)

Transmission Control Protocol, Src Port: 1065 (1065), Dst Port: http (80), Seq: 1, Ack: 1, Len: 428

Hypertext Transfer Protocol

GET /~knuth/ HTTP/1.1\r\n

Host: www-cs-faculty.stanford.edu\r\n

User-Agent: Mozilla/5.0 (Windows; U; Windows NT 5.1; en-GB; rv:1.8.1.12) Gecko/20080201 Firefox/2.0.0.12\r\n

Accept: text/xml,application/xml,application/xhtml+xml,text/html;q=0.9,text/plain;q=0.8,image/png,*/*;q=0.5\r\n

Accept-Language: en-gb,en;q=0.5\r\n

Accept-Encoding: gzip,deflate\r\n

Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7\r\n

Keep-Alive: 300\r\n

Connection: keep-alive\r\n

\r\n

File: "C:\DOCUMENT~1\Kevin\LOCALS~1\Temp\etherXXXXa02592" 47 KB 00:00:10

P: 122 D: 122 M: 0 Drops: 0

Start (Untitled) - Wireshark Don Knuth's Home Page ... Jasc Paint Shop Pro 08:54

The HTTP Response

The server sends back something like:

HTTP/1.1 200 OK

Date: Mon, 23 May 2005 22:38:34 GMT

Server: Apache/1.3.27 (Unix) (Red-Hat/Linux)

Last-Modified: Wed, 08 Jan 2003 23:11:55 GMT

Content-Length: 438

Connection: close

Content-Type: text/html; charset=UTF-8

<html>

<body>

and the rest of the requested resource

HTTP Request Format

Request Line:

`GET /index.html HTTP/1.1`

Method URI Version No

Followed by a set of headers:

`Host: www.example.com`

`User-Agent: Mozilla/5.0 (Linux; X11; UTF-8)`

`Accept-Language: en`

Sometimes followed by an optional message body

HTTP Request Methods

GET

- Requests a copy of the resource stored at the specified URI.

POST

- Submits data to be processed (e.g. from an HTML form) to the specified URI. The data is included in the message body.

PUT

- Requests the server to store the data enclosed in the message body at the specified URI.

Some of the less common HTTP request methods are shown next.

Be aware that not all web servers implement all aspects of the HTTP protocol.

HTTP Request Methods

HEAD

- Same as a GET request, but tells the server to just send the response headers, not the message body.

DELETE

- Requests the server to delete the resource stored at the specified URI.

TRACE

- Requests that the server echoes back the received request, so that a client can see what intermediate servers are adding or changing in the request. Used for testing.

OPTIONS

- Returns the HTTP methods that the server supports. This can be used to check the functionality of a web server.

CONNECT

- Requests the server to set up an secure (usually encrypted) connection.

Safe Methods

GET and HEAD are conventionally regarded as safe methods.

They should only be used for data retrieval and should not be used to change the state of the server.

POST, PUT and DELETE can be used to change the state of the server.

Idempotent Methods

GET, HEAD, PUT and DELETE are defined as idempotent methods.

ie Multiple identical requests should have the same effect as a single request.

This allows a user agent to retry idempotent requests without informing the user, which may be useful when connecting to busy servers.

Which makes it more important to not use GET for anything other than document retrieval.

HTTP Response Format

Status Line:

HTTP/1.1 200 OK

| Version | Status Code | Reason Phrase |
|---------|-------------|---------------|
|---------|-------------|---------------|

Followed by a set of headers

```
Date: Mon, 23 May 2005 22:38:34 GMT
Server: Apache/1.3.27 (Unix) (Red-Hat/Linux)
Last-Modified: Wed, 08 Jan 2003 23:11:55 GMT
Content-Length: 438
Connection: close
Content-Type: text/html; charset=UTF-8
```

and a message body, sometimes called an entity

```
<html>
<body>
and so on
```

HTTP Response Status Codes

1xx Informational

100 Continue

101 Switching Protocols

2xx Success

200 OK

201 Created

202 Accepted

203 Non-Authoritative Information

204 No Content

205 Reset Content

206 Partial Content

HTTP Response Status Codes

3xx Redirection

The client must take additional action to complete the request. This may be done by the user agent without notifying the user.

| | |
|-----|--------------------|
| 300 | Multiple Choices |
| 301 | Moved Permanently |
| 302 | Found |
| 303 | See other |
| 304 | Not Modified |
| 305 | Use Proxy |
| 307 | Temporary Redirect |

HTTP Response Status Codes

4xx Client Error

The request contains bad syntax or cannot be fulfilled. The server response should include a message entity containing an explanation of the situation.

| | |
|-----|----------------------|
| 400 | Bad Request |
| 401 | Unauthorised |
| 402 | Payment Required |
| 403 | Forbidden |
| 404 | Not Found |
| 405 | Method Not Allowed |
| 406 | Not Acceptable |
| 408 | Request Timeout |
| 414 | Request URI Too Long |

HTTP Response Status Codes

5xx Server Error

The server failed to fulfill an apparently valid request. The server response should contain a message entity containing an explanation of the error situation.

| | |
|-----|---------------------------------|
| 500 | Internal Server Error |
| 501 | Not Implemented |
| 502 | Bad Gateway |
| 503 | Service Temporarily Unavailable |
| 504 | Gateway Timeout |
| 505 | HTTP Version Not Supported |

Persistent Connections

In HTTP/0.9 and 1.0, the connection is closed after a single request/response pair.

In HTTP/1.1 a keep-alive-mechanism was introduced, where a connection could be reused for more than one request.

Such *persistent connections* reduce lag perceptibly, because the client does not need to re-negotiate the TCP connection after the first request has been sent.

HTTP Session State

HTTP is a stateless protocol.

Each request is treated as an independent transaction, unrelated to any other request.

This forces the use of alternative methods for maintaining a users' state, for example, when a host would like to customize content for a user who has visited before.

The common methods for solving this problem are:

- sending and requesting cookies.
- server side sessions.
- hidden variables (when current page is a form).
- URL encoded parameters (such as `/index.php?userid=3`).